Meeting Follow-Up | December 2, 2025

Griffith Foods + nBrain

Introductory Call Summary & Your Questions Answered

A comprehensive recap of our discussion, with detailed answers to every question raised by your team, and a focused look at how The Griffith Brain can transform Product Development.

Meeting Summary

DATE

Tuesday, December 2, 2025

GRIFFITH FOODS

TC (CEO), Mike Cozzi (Global IT), Chris Evans (Data Architecture), Matt Corker (CFO)

NBRAIN TEAM

Danny DeMichele, Dan Caufield, Patrick Laughlin

PRIMARY FOCUS

Product Development AI Transformation

Key Themes Discussed

The Brain Architecture

Your private, agentic AI system that learns from every interaction and compounds intelligence over time. Not a generic AI tool - a proprietary intelligence layer built for Griffith.

Work-For-Hire Model

Everything we build belongs to you. Full source code, documentation, and capability transfer. Our goal is to replace ourselves by building your internal AI team.

Change Management First

AI transformation fails without people buy-in. We use Kotter's methodology and focus on building champions, not just deploying technology.

Your Strategic Vision (As We Understand It)

Griffith Foods operates across the value chain: **Source - Develop - Make - Deliver**.

You've identified the **DEVELOP** phase (Product Development/R&D) as the starting point for AI transformation because that's where you create the most value for customers. The goal: implement AI properly in DEVELOP, then scale the learnings across your entire value chain.

The BHAG: Become the leading AI-driven partner for product development in the food industry - and grow from $2 billion to $4 billion.

Your Questions, Answered

We captured every question from your team during our conversation. Below are comprehensive answers to each one, with the detail and clarity you deserve.

TC (CEO)

"Tell us a little bit more about the company itself - how you're organized, how many people, types of clients you have."

nBrain at a Glance

Core Team: ~10 full-time employees, augmented by 25-30 specialized contractors (1099s), with an additional 50 vetted professionals on our bench ready to be deployed for specific projects.

Why We're Small by Design: We drink our own Kool-Aid. Every one of our employees has 10-15x productivity because we built our company AI-first from day one. We spent the first two years building our own internal AI infrastructure before bringing on clients. This means a team of 10 can outproduce what traditionally takes 50+ people.

Client Profile: Primarily mid-market companies ($25M - $100M revenue), though we've worked with companies ranging from $2M to multi-billion dollar enterprises. Our sweet spot is companies that:

- Have valuable proprietary data and processes worth protecting
- Want to build internal AI capability, not just buy tools
- Are committed to change management, not just technology deployment
- Have leadership buy-in for AI transformation

Why Griffith Foods Is Ideal For Us: You represent everything we look for - real proprietary secrets that must stay secret (your 100 years of formulation knowledge), a clear strategic vision (DEVELOP phase first), executive commitment, and the scale to truly benefit from compounding AI intelligence.

Chris Evans

"What's your initial approach to a company like ours? Give me an anchor point - where does this conversation start? How do you help us begin the AI journey?"

Our Engagement Starts With Three Things

1. Data

Everything starts with understanding your data. Not just what systems you have, but what data is AI-ready, what needs preparation, and what represents your highest-value opportunities.

2. People

We identify your champions - the 15% who are curious and eager, not resistant. We also map potential resistance points so we can address them proactively.

3. Quick Wins

We identify 2-3 use cases that can demonstrate value quickly, building momentum and buy-in for larger transformation.

The 90-Day Foundation Engagement

Month 1: Dual Audit

Data + People Discovery

We audit your technology AND your organizational readiness. This means stakeholder interviews at the vision, management, and operational levels. We map current workflows, identify data sources, and assess AI readiness - not just technically, but culturally.

Month 2: Architecture Build

Build The Griffith Brain Foundation

We don't just design a blueprint - we build the infrastructure on YOUR stack. This includes your private vector database,

the master agent process, security framework, and data pipelines. Your team participates in every step.

Month 3: Use Case Validation

Decision Lab Process

We put potential projects through our Decision Lab - a rigorous validation process that ensures we're not pursuing "willy-nilly" projects. You'll have 3 validated use cases ready to build, complete with data mapping, KPIs, and success criteria.

KEY POINT

You build it, we provide the blueprints and expertise. Everything is done on your stack, with your people, so there's no dependency on us. At any point, you can take everything and run with it independently.

Mike Cozzi

"What does the project team look like at the beginning? What does it evolve to over time? What's the skill set and resource type we need to make this successful?"

Our Team Dedicated to You

Executive Sponsor

Dan Caufield or Danny DeMichele - Strategic alignment with your steering committee

Program Manager

Patrick Laughlin - Day-to-day execution, coordinates with your Tiger Team

AI Architect

Designs Brain architecture, security framework, integration points

Data Scientist

Data audit, readiness assessment, vectorization strategy

AI Engineers (2)

Build agents, implement workflows, work alongside your developers

What We Need From Griffith Foods

At the beginning: You don't need to hire anyone new. We need:

- **An Executive Sponsor:** Someone on your steering committee who champions this initiative
- **Tiger Team Access:** Cross-functional leaders who can make decisions and provide access to stakeholders

- **Interview Time:** Access to key people in your DEVELOP phase for discovery conversations

How It Evolves

Over time, your team grows organically. The people we work with become your AI champions. They don't need AI backgrounds - they need:

- **Curiosity over credentials:** The best AI team members are curious, not resistant. They may be using ChatGPT casually, but what matters is their desire to learn.
- **Domain expertise:** Your product development scientists know your processes better than any AI expert ever will. That domain knowledge is more valuable than AI skills.
- **Management ability:** They need to understand workflows and processes, not write code.
- **Staying power:** Ideally, these are people you expect to be with Griffith long-term.

THE KEY INSIGHT

AI does the AI. When you have your own private Brain, the AI itself knows everything about AI. Your people don't need to be prompting experts - the agents are pre-trained with context, rules, tone, and history. Users just ask questions in natural language. The complexity is handled by the architecture, not by training every employee.

Mike Cozzi "Is the Brain the same as the LLM? When you say we'll have our own LLM, is that what's built for us?"

Two Different Things - Both Important

The Griffith Brain (Immediate)

What it is: A private instance, vectorized, agentic process completely trained on YOUR processes, systems, and data.

When you get it: Built during the 90-day engagement and immediately usable.

How it works: Uses existing large language models (OpenAI, Claude, etc.) in a secure, multimodal configuration. The Brain decides which model to use for each task. All your data stays within your infrastructure - the LLMs never train on your information.

What it does: Allows you to use AI at a very high level, in a completely secure environment, with agents pre-configured for your specific workflows.

Your Micro-LLM (2-3 Years)

What it is: A proprietary, domain-specific small language model trained entirely on Griffith Foods' data and interactions.

When you get it: Typically 2-3 years. It requires significant accumulated training data from Brain interactions.

How it works: The "exhaust" from every Brain interaction - every query, every agent task, every human feedback loop - becomes training data for your own model.

Why it matters: At this point, you may not need external LLMs at all. Your Micro-LLM encodes 100 years of Griffith expertise plus years of AI-augmented learning. It becomes a strategic asset potentially worth more than physical assets.

Think of it This Way

The Brain is the vehicle. The Micro-LLM is the destination.

The Brain gives you immediate value - productivity, intelligence, and the ability to leverage AI across your organization TODAY. But every interaction with the Brain is also generating the training data you'll eventually use to create something no competitor can replicate: your own Griffith-specific AI model.

This is the "compounding intelligence" we talked about. Every day, the Brain gets smarter. Every day, you're accumulating the data for something even more powerful.

Mike Cozzi

"Once we put the Brain in place, what skill set will people need to leverage it? There has to be some kind of skillset to use the brain, I'm assuming."

The Short Answer: Less Than You Think

> **If you do AI right, adopting it shouldn't require training your staff.**
>
> The problems you see with ChatGPT - needing to be a "prompting genius," inconsistent results, hallucinations - those go away with your own private Brain. Why? Because we train the AI at the agent level. Context is already provided. Rules, tone, history - everything is locked in.

What "Using the Brain" Actually Looks Like

For Your Product Development Scientists

They ask questions in natural language. "What similar formulations have we done for clean-label seasoning blends?" or "What are the FDA requirements for this ingredient combination?" The Brain handles the complexity - querying your 100 years of data, checking regulations, finding relevant precedents.

Required training: Maybe 30 minutes to show them how to access it. The interface can be embedded in tools they already use - their PLM system, email, Slack, even text.

For Your Champions (Task Force Members)

These are the people who will help identify new use cases and work with us to validate them. They need to understand the art of the possible, not the technical details. They'll participate in our Decision Lab process and learn to think "AI-first" about problems.

Required training: Ongoing as they work with our team. This is embedded learning, not classroom training.

The Real Skill: Curiosity, Not Credentials

The people who succeed with AI aren't the ones with technical backgrounds. They're the ones who:

- Ask "why not?" instead of "why"
- Are willing to experiment and provide feedback

- Understand their domain deeply (your product development processes)
- Can articulate what problems they want solved

THE COUNTERINTUITIVE TRUTH

Learning how to prompt is great if you're a 1-2 person shop. At a company of your size, you'll never get AI adopted by teaching everyone to use AI in the generic sense. Instead, we embed AI into the systems they already use, pre-configure it for their specific workflows, and make it proactive where possible. They don't need to know how to use AI - they just need to do their jobs, and AI helps them do it better.

Chris Evans

"Change management becomes huge - that's a concern. Also, is there a function of the brain that interviews people? Does it have the ability to ask questions, not just answer them?"

Addressing Change Management

You're right - this is huge, and it's why most AI implementations fail. The technology is the easy part. Getting people to adopt it is the hard part.

Our Change Management Approach

Based on Patrick's experience with 200-300 companies, here's the reality:

- **~40% of your workforce will be resistant** to any AI automation. That's normal. We plan for it.
- **~15% will be early adopters** who are excited and may even go too far. These become your champions.
- **The remaining 45%** will follow whichever group has momentum. Our job is to create momentum with wins.

We follow Kotter's 8-step change model: Create urgency, build coalition, form vision, enlist volunteers, enable action, generate wins, sustain acceleration, institute change.

Can the Brain Be Proactive? (Ask Questions, Not Just Answer)

Yes, absolutely. True AGI (Artificial General Intelligence) isn't here yet - AI isn't truly proactive by nature. But we create proactive behaviors through what we call "artificial hats":

Scheduled Intelligence

Like old-school cron jobs, but smart. We can tell the Brain: "Every day, run analysis on all new formulation tests and flag

anomalies we should be concerned about."

Or: "Every time a customer submits a new product brief, automatically check it against our existing formulations and suggest starting points."

Triggered Actions

Web hooks and events can trigger AI analysis. When data enters the system - say, new quality test results - the Brain automatically performs pre-structured analysis and surfaces insights.

When someone's working on a formulation, the Brain can automatically check for regulatory conflicts or flag potential issues before they become problems.

OUR GOAL

If AI is done right, it's 75%+ proactive and only 25% reactive. We don't want people to have to remember to ask. We want AI surfacing insights, flagging issues, and anticipating needs before humans even know to ask.

Key Discussion

"How are you currently capturing the collective intelligence of your best people?"

This question - asked by Patrick to your team - gets to the heart of what makes the Brain valuable.

The Core Problem

Your best product development scientists have years, sometimes decades, of expertise. But right now, that knowledge lives in their heads, their notebooks, or scattered documents. When they leave, retire, or move to new roles, that knowledge goes with them.

The Brain captures that collective intelligence in every interaction.

How It Works

- Every question asked, every formulation attempted, every successful outcome - it all becomes part of the Brain's knowledge base
- You can curate what it learns: "Yes, learn from this successful formulation" or "No, don't train on that failed experiment"
- Over time, the Brain contains not just what's in your databases, but the accumulated wisdom of your best people
- New scientists can tap into decades of expertise immediately

THE COMPETITIVE MOAT

This is your Coca-Cola recipe in a vault. Collective intelligence that is incredibly valuable, proprietary, and impossible for competitors to replicate. Combined with the billions of dollars of investment in the underlying LLMs, you get the best of both worlds: cutting-edge AI capability PLUS your unique, irreplaceable domain expertise.

Product Development: Where AI Creates Maximum Impact

You've identified the DEVELOP phase as your starting point for AI transformation. Based on what we understand about Griffith Foods' business - over 100 years of expertise, 27 global manufacturing sites, and product categories spanning seasonings, sauces, coatings, functional blends, and more - here's how The Griffith Brain could transform your Product Development process.

Griffith Foods' Product Development Strength

From griffithfoods.com: *"Griffith Foods is the intersection of culinary arts, leading-edge restaurant and food knowledge, and proven food science. At Griffith Foods, we have perfected the integration of culinary arts, food science, sensory science and consumer insights to deliver best-in-class products and solutions."*

The opportunity: What if AI could help your scientists leverage that integrated expertise - culinary arts, food science, sensory science, and consumer insights - all at once, for every formulation decision?

The Current Product Development Challenge

Data Is Scattered

- Formulation history across 27 sites
- Quality test results in different systems
- Regulatory requirements that vary by region
- Customer feedback in CRM and emails
- Sensory test results in spreadsheets

Knowledge Lives in Heads

- Best practices aren't consistently documented
- When experts retire, knowledge leaves
- Regional teams may duplicate work
- New scientists have steep learning curves
- Tribal knowledge is hard to transfer

How The Griffith Brain Changes This

1

Customer Brief

2

Formulation

3

Testing

Intake & Analysis AI-Assisted R&D Quality & Sensory

4

Validation

Compliance Check

5

Scale-Up

Manufacturing

Potential AI-Enhanced Product Development Use Cases

Note: These are illustrative examples. Actual use cases will be identified and validated during discovery with your team.

1 Formulation Intelligence Assistant

> **Scenario:** A product development scientist receives a brief from a QSR customer: "We need a clean-label seasoning blend for our new plant-based burger that delivers an umami flavor profile, holds up during high-heat cooking, and costs less than $X per pound. Target launch: 90 days."

How the Brain Helps:

- **Instant precedent search:** "Show me all clean-label seasonings we've developed for plant-based applications in the last 3 years" - across all 27 sites
- **Ingredient intelligence:** "What natural umami boosters have we used successfully? Which ones are cost-effective?"
- **Regulatory check:** "Are there any allergen or labeling concerns for this ingredient combination in the US and EU?"

- **Cost modeling:** "Based on current commodity prices, what's the estimated cost per pound for this formulation?"
- **Quality prediction:** "Based on similar formulations, what quality attributes should we monitor during testing?"

Impact: What might take days of searching databases and consulting colleagues could happen in minutes, with the collective wisdom of your entire organization available instantly.

2 Cross-Regional Knowledge Sharing

> **Scenario:** Your European team is developing a new textures & coatings product for a customer's fish application. Meanwhile, your North American team solved a similar challenge last year for poultry.

How the Brain Helps:

- **Automatic connection:** When the EU team starts working on the fish coating, the Brain surfaces: "FYI: The NA team developed a similar high-adhesion coating for poultry in Q3 2024. Key learnings attached."
- **Knowledge transfer:** Best practices, failed approaches, and quality test results from the NA project are immediately accessible
- **Adaptation suggestions:** "Based on fish vs. poultry surface characteristics, here are recommended adjustments to the base formulation."

Impact: Eliminate duplicated R&D work across regions. Leverage the full breadth of Griffith's global expertise for every project.

3 Regulatory Compliance Automation

> **Scenario:** A functional blend is being developed for a customer who sells in the US, EU, and Asia-Pacific markets. Each region has different regulatory requirements for the same ingredients.

How the Brain Helps:

- **Multi-jurisdictional check:** Automatically validates ingredient combinations against FDA, EU food safety, and regional APAC regulations
- **Allergen tracking:** Identifies all allergen declaration requirements by market
- **Label verification:** Ensures ingredient naming conventions meet each region's requirements
- **Change monitoring:** When regulations change (which they do constantly), the Brain alerts relevant projects that may be affected

Impact: Reduce compliance review time from days to hours. Catch potential issues before they become costly reformulation requirements.

4 Consumer Trend-to-Formulation Pipeline

> **Scenario:** Consumer data shows rising interest in "mushroom-based flavors" and "adaptogens." Your sales team wants to proactively present innovation concepts to key customers.

How the Brain Helps:

- **Trend analysis:** "What consumer trends align with our current capabilities in seasonings and functional blends?"
- **Formulation concepts:** "Generate 3 concept directions for mushroom-forward seasonings based on our existing ingredient portfolio"
- **Customer matching:** "Which of our top 20 customers have expressed interest in functional ingredients or mushroom flavors?"
- **Competitive intelligence:** "What mushroom-based products have competitors launched in the last 12 months?"

Impact: Transform consumer insights into actionable product concepts faster. Be proactive with customers rather than reactive to briefs.

AI-Enhanced Product Categories

Based on your product portfolio from griffithfoods.com, here's how AI could enhance each category:

Seasonings

Rubs, marinades, glazes, internal seasonings

AI ENHANCEMENT

Flavor profile matching, regional taste preference analysis, cost optimization across ingredient alternatives

Sauces & Dressings

New and exciting blends for any liquid application

AI ENHANCEMENT

Viscosity prediction, shelf-life optimization, clean-label reformulation suggestions

Textures & Coatings

Cereal and flour-based coatings with added flavors

AI ENHANCEMENT

Adhesion optimization, crispness retention modeling, application method recommendations

Functional Blends

High-performance ingredient blends for meat, poultry, seafood

AI ENHANCEMENT

Yield prediction, cook loss optimization, moisture retention analysis

Soups & Sides

Innovative side options with global flavors

AI ENHANCEMENT

Global flavor trending, sodium reduction alternatives, reconstitution optimization

Dough Blends

Bakery and dough systems for taste, texture, performance

AI ENHANCEMENT

Fermentation optimization, texture profiling, freeze-thaw stability prediction

What AI Already Knows About Product Development

As Danny mentioned in our call: there's a lot that AI knows out of the box. You almost don't need to train it for certain things.

Base Knowledge Already Available

- **FDA guidelines** around ingredients, labeling requirements, GRAS status
- **EU food safety standards** and regulatory frameworks across European markets
- **Ingredient interactions** - how different compounds affect each other chemically and functionally
- **Food science fundamentals** - emulsification, protein functionality, starch behavior, etc.
- **Global regulatory variations** - how requirements differ by country/region
- **Consumer trend data** - plant-based, clean label, functional ingredients, etc.
- **Competitive landscape** - publicly available product launches and innovations

Your Layer on Top

The base knowledge is powerful, but it's generic - your competitors could access the same LLMs. What makes your Brain unique is what you layer on top:

- **100 years of Griffith formulation history** - successful approaches, failed experiments, lessons learned
- **Your specific quality standards** - what "good" looks like for Griffith
- **Customer relationships and preferences** - what each major customer values
- **Internal processes and workflows** - how Griffith actually develops products
- **Pricing and cost structures** - real economics, not industry averages
- **Your people's expertise** - captured in every interaction with the Brain

This combination - base AI intelligence + proprietary Griffith knowledge - is what creates the compounding competitive advantage.

The Compounding Intelligence Advantage

How It Works

Every single interaction with the Brain generates data:

- Every question asked by your scientists
- Every formulation attempt and outcome
- Every quality test result
- Every customer feedback point
- Every decision made

This creates **synthetic data** - insights that never existed before. Each interaction makes the Brain smarter, and each improvement compounds.

Why It's a Moat

Day 1: Your Brain knows what you teach it - your historical formulations, processes, standards.

Day 365: It knows everything from Day 1, PLUS a year of accumulated interactions, insights, and learnings.

Day 1,000: It contains nearly three years of compounded intelligence - instantly accessible, continuously improving.

A competitor starting today can never catch up. They'd have to start from scratch while you're accelerating away.

THE HARVARD STUDY DAN REFERENCED

Researchers gave an AI nothing but data points about ball positions at different times. From just that raw data, the AI derived velocity. From velocity, it derived mass. Then gravity. Then all the laws of physics - through millions of iterations and pattern recognition. **That's what happens when AI has access to enough of the right data.** Your 100 years of formulation history, combined with ongoing operational data, gives you that same opportunity for your domain.

Protecting Your Secret Recipes

As Dan said: "You have real secrets. Those secrets cannot be let out."

What You're Protecting

- 100 years of proprietary formulations
- Customer relationships and specific needs
- Pricing and cost structures
- Quality standards and processes
- Competitive intelligence
- Your people's accumulated expertise

How We Protect It

- **Private infrastructure:** Everything runs on YOUR stack
- **Zero external training:** LLMs don't learn from your data
- **Tokenized processing:** Data is encrypted before AI processes it
- **Role-based access:** Different visibility for different roles
- **Complete audit trails:** Every query logged for compliance

- **LLM-agnostic:** Swap models without changing architecture

The Difference From ChatGPT/Claude Direct Use

When your team uses ChatGPT or Claude directly with proprietary data, you're potentially exposing competitive information. Those platforms may use your inputs to train future models. Even if they claim they don't, you have no control.

With The Griffith Brain, **you control everything**. The LLMs are used as reasoning engines only - they never see your actual data in readable form. All the knowledge lives in your private vector database, encrypted and isolated.

Proposed Next Steps

Step 1: Internal Discussion

Review This Document as a Team

Share this follow-up with your Tiger Team. Discuss which questions were answered, which need more clarity, and what concerns remain.

Step 2: Follow-Up Call

Deeper Dive on Specific Areas

We're happy to schedule another call to go deeper on any topics - whether that's technical architecture, change management approach, product development use cases, or economic model.

Step 3: Meet the Team

Introduction to Specific People

If you want to proceed further, we'll introduce you to the specific team members who would work on your account - so you can evaluate not just our company, but the actual people.

Step 4: Statement of Work

Formalize the 90-Day Engagement

If there's mutual fit, we draft a detailed SOW covering scope, deliverables, timeline, and economics for the 90-day foundation engagement.

We're Ready When You Are

This isn't a hard sell. We want you to be confident this is the right fit for Griffith Foods. Take your time, ask us hard questions, and let's figure out together if this makes sense.

Contact: Danny DeMichele | Dan Caufield | Patrick Laughlin

Questions We Didn't Get To

We ran out of time, but there are likely questions your team still has. Here are some common ones we're happy to address:

Economic Model Details

The 90-day engagement, monthly retainer structure, the "happiness bonus" concept, and how costs scale as you grow.

Technical Deep Dive

Vector database specifics, RAG architecture, integration with your PLM/ERP systems, security certifications.

Case Studies

While we haven't worked in food specifically, we can share learnings from similar IP-sensitive industries and change management successes.

Timeline to Value

When can you expect to see ROI? What does the first "win" typically look like? How do we measure success?

nBrain Platform

Proprietary AI Intelligence Systems for Enterprise

Building AI that compounds. Building teams that last.